

Artificial Intelligence

For HEDSPI Project

Lecturer 5 - Advanced search methods

Lecturers :

Dr. Le Thanh Huong

Dr. Tran Duc Khanh

Dr. Hai V. Pham

HUST

1

Outline

- Memory-bounded heuristic search
- Hill-climbing search
- Simulated annealing search

2
2

Memory-bounded heuristic search

- Some solutions to A* space problems (maintain completeness and optimality)
 - Iterative-deepening A* (IDA*)
 - Here cutoff information is the f -cost ($g+h$) instead of depth
 - Recursive best-first search(RBFS)
 - Recursive algorithm that attempts to mimic standard best-first search with linear space.
 - (simple) Memory-bounded A* ((S)MA*)
 - Drop the worst-leaf node when memory is full

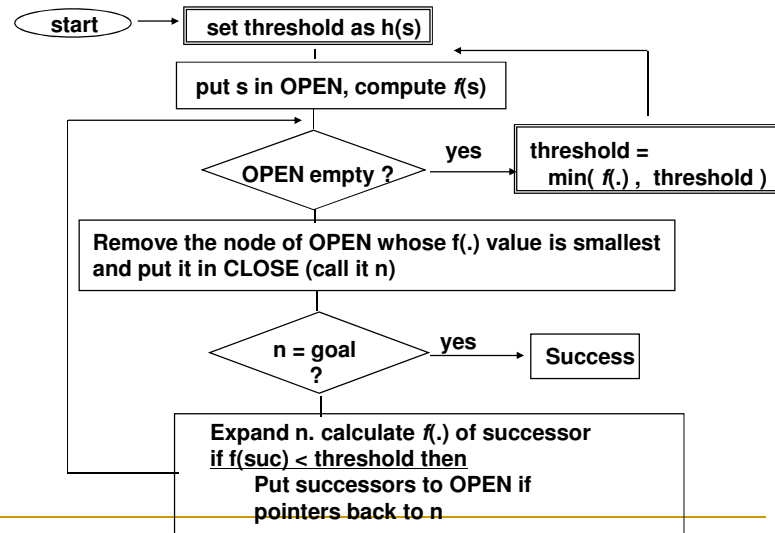
3

Iterative Deeping A*

- Iterative Deeping version of A*
 - use threshold as depth bound
 - To find solution under the threshold of $f(.)$
 - increase threshold as minimum of $f(.)$ of
 - previous cycle
- Still admissible
- same order of node expansion
- Storage Efficient – practical
 - but suffers for the real-valued $f(.)$
 - large number of iterations

4

Iterative Deepening A* Search Algorithm (for tree search)



5

Recursive best-first search

- A variation of Depth-first search
- Keep track of f -value of the best alternative path
- Unwind if f -value of all children exceed its best alternative
- When unwind, store f -value of best child as its f -value
- When needed, the parent regenerate its children again.

6

Recursive best-first search

function RECURSIVE-BEST-FIRST-SEARCH(*problem*) **return** a solution or failure
 return RBFS(*problem*, MAKE-NODE(INITIAL-STATE[*problem*]), ∞)

function RBFS (*problem*, *node*, *f_limit*) **return** a solution or failure and a new *f-cost*
 limit
 if GOAL-TEST[*problem*](STATE[*node*]) **then return** *node*
 successors \leftarrow EXPAND(*node*, *problem*)
 if *successors* is empty **then return** failure, ∞
 for each *s* **in** *successors* **do**
 $f[s] \leftarrow \max(g(s) + h(s), f[node])$
 repeat
 best \leftarrow the lowest *f*-value node in *successors*
 if $f[best] > f_limit$ **then return** failure, $f[best]$
 alternative \leftarrow the second lowest *f*-value among *successors*
 result, $f[best] \leftarrow$ RBFS(*problem*, *best*, $\min(f_limit, alternative)$)
 if *result* \neq failure **then return** *result*

7


Recursive best-first search

- Keeps track of the *f*-value of the best-alternative path available.
 - If current *f*-values exceeds this alternative *f*-value then backtrack to alternative path.
 - Upon backtracking change *f*-value to best *f*-value of its children.
 - Re-expansion of this result is thus still possible.

8

Recursive best-first search, ex.

(a) After expanding Arad, Sibiu, Rimnicu Vilcea




- Path until Rimnicu Vilcea is already expanded
- Above node; f -limit for every recursive call is shown on top.
- Below node: $f(n)$
- The path is followed until Pitesti which has a f -value worse than the f -limit.

9

Recursive best-first search, ex.

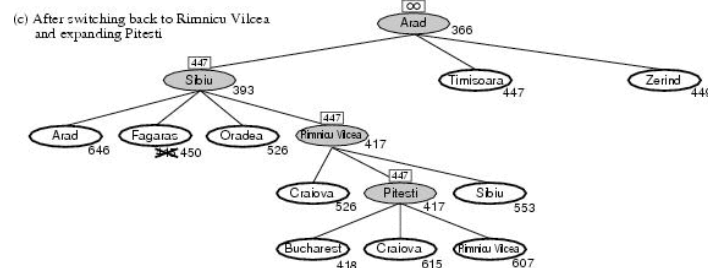
(b) After unwinding back to Sibiu and expanding Fagaras



- Unwind recursion and store best f -value for current best leaf Pitesti
 $result, f[best] \leftarrow RBFS(problem, best, \min(f_limit, alternative))$
- $best$ is now Fagaras. Call RBFS for new $best$
 - $best$ value is now 450

10

Recursive best-first search, ex.



- Unwind recursion and store best f -value for current best leaf Fagaras
 $result, f[best] \leftarrow RBFS(problem, best, \min(f_limit, alternative))$
- $best$ is now Rimnicu Vilcea (again). Call RBFS for new $best$
 - Subtree is again expanded.
 - Best *alternative* subtree is now through Timisoara.
- Solution is found since because $447 > 417$.

11

RBFS evaluation

- RBFS is a bit more efficient than IDA*
 - Still excessive node generation (mind changes)
- Like A*, optimal if $h(n)$ is admissible
- Space complexity is $O(bd)$.
 - IDA* retains only one single number (the current f -cost limit)
- Time complexity difficult to characterize
 - Depends on accuracy of $h(n)$ and how often best path changes.
- IDA* and RBFS suffer from **too little** memory.

12

(simplified) memory-bounded A*

- Use all available memory.
 - I.e. expand best leafs until available memory is full
 - When full, SMA* drops worst leaf node (highest f -value)
 - Like RBFS, we remember the best descendant in the branch we delete
- What if all leafs have the same f -value?
 - Same node could be selected for expansion and deletion.
 - SMA* solves this by expanding *newest* best leaf and deleting *oldest* worst leaf.
- The deleted node is regenerated when all other candidates look worse than the node.
- SMA* is complete if solution is reachable, optimal if optimal solution is reachable.
- Time can still be exponential.


13

Local search algorithms

- In many optimization problems, the **path** to the goal is irrelevant; the goal state itself is the solution
- State space = set of "complete" configurations
- Find configuration satisfying constraints, e.g., n-queens
- In such cases, we can use **local search algorithms**
- Local search= use single current state and move to neighboring states.
- Advantages:
 - Use very little memory
 - Find often reasonable solutions in large or infinite state spaces.
- Are also useful for pure optimization problems.
 - Find best state according to some *objective function*.

14


Local search and optimization



15

Example: n -queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal



16

Hill-climbing search

- Simple, general idea:
 - Start wherever
 - Always choose the best neighbor
 - If no neighbors have better scores than current, quit
- Hill climbing does not look ahead of the immediate neighbors of the current state.
- Hill-climbing chooses randomly among the set of best successors, if there is more than one.
- Some problem spaces are great for hill climbing and others are terrible.

17

Hill-climbing search

function HILL-CLIMBING(*problem*) **return** a state that is a local maximum

input: *problem*, a problem

local variables: *current*, a node.

neighbor, a node.

current ← MAKE-NODE(INITIAL-STATE[*problem*])

loop do


neighbor ← a highest valued successor of *current*

if VALUE [*neighbor*] < VALUE[*current*] **then return** STATE[*current*]

current ← *neighbor*

18


Robot Navigation



$$f(N) = h(N) = \text{straight distance to the goal}$$

19


Hill climbing example



$$f(n) = -(\text{number of tiles out of place})$$

20

Example of a local maximum




21

Drawbacks of hill climbing

- Problems:
 - Local Maxima:** depending on initial state, can get stuck in local maxima
 - Plateaus:** the space has a broad flat region that gives the search algorithm no direction (random walk)
 - Ridges:** flat like a plateau, but with dropoffs to the sides; steps to the North, East, South and West may go down, but a combination of two steps (e.g. N, W) may go up

➤ Introduce randomness



Hill-climbing variations

- Stochastic hill-climbing
 - Random selection among the uphill moves.
 - The selection probability can vary with the steepness of the uphill move.
- First-choice hill-climbing
 - Stochastic hill climbing by generating successors randomly until a better one is found.
- Random-restart hill-climbing
 - Tries to avoid getting stuck in local maxima.
 - If at first you don't succeed, try, try again...


23

Simulated Annealing

- Simulates slow cooling of annealing process
- Applied for combinatorial optimization problem by S. Kirkpatrick ('83)
- **What is annealing?**
 - Process of slowly cooling down a compound or a substance
 - Slow cooling let the substance flow around → thermodynamic equilibrium
 - Molecules get optimum conformation

24

Simulated annealing



gradually decrease shaking to make sure the ball escape from local minima and fall into the global minimum

25

Simulated annealing

- Escape local maxima by allowing “bad” moves.
 - Idea: but **gradually decrease** their size and frequency.
- Origin; metallurgical annealing
- Implement:
 - Randomly select a move instead of selecting best move
 - Accept a bad move with probability less than 1 ($p < 1$)
 - p decreases by time
- If T decreases slowly enough, best state is reached.
- Applied for VLSI layout, airline scheduling, etc.

26

Simulated annealing

function SIMULATED-ANNEALING(*problem*, *schedule*) **return** a solution state

input: *problem*, a problem

schedule, a mapping from time to temperature

local variables: *current*, a node; *next*, a node.

T, a “temperature” controlling the probability of downward steps

current \leftarrow MAKE-NODE(INITIAL-STATE[*problem*])

for *t* \leftarrow 1 **to** ∞ **do**

T \leftarrow *schedule*[*t*]

if *T* = 0 **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ VALUE[*next*] - VALUE[*current*]

if $\Delta E > 0$ **then** *current* \leftarrow *next*

else *current* \leftarrow *next* only with probability $e^{\Delta E / T}$

Similar to hill climbing,
but a **random** move
instead of best move

case of improvement, make the move

Otherwise, choose the move with
probability that decreases exponentially
with the “badness” of the move.

What’s the probability when: $T \rightarrow \text{inf}$?

What’s the probability when: $T \rightarrow 0$?

What’s the probability when: $\Delta = 0$?

What’s the probability when: $\Delta \rightarrow -\infty$?

Simulated Annealing parameters

■ Temperature *T*


- Used to determine the probability
- High *T* : large changes
- Low *T* : small changes

■ Cooling Schedule

- Determines rate at which the temperature *T* is lowered
- Lowers *T* slowly enough, the algorithm will find a global optimum

- In the beginning, aggressive for searching alternatives, become conservative when time goes by

Simulated Annealing Cooling Schedule



- if T_f is reduced too fast, poor quality
- if $T_t \geq T(0) / \log(1+t)$ - Geman
 - System will converge to minimum configuration
- $T_t = k/1+t$ - Szu
- $T_t = a \cdot T(t-1)$ where a is in between 0.8 and 0.99

29

Tips for Simulated Annealing

- To avoid of entrainment in local minima
 - Annealing schedule : by trial and error
 - Choice of initial temperature
 - How many iterations are performed at each temperature
 - How much the temperature is decremented at each step as cooling proceeds
- Difficulties
 - Determination of parameters
 - If cooling is too slow → Too much time to get solution
 - If cooling is too rapid → Solution may not be the global optimum

30

Properties of simulated annealing

- Theoretical guarantee:
 - Stationary distribution: $p(x) \propto e^{-\frac{E(x)}{kT}}$
 - If T decreased slowly enough, will converge to optimal state!
- Is this an interesting guarantee?
- Sounds like magic, but :
 - The more downhill steps you need to escape, the less likely you are to every make them all in a row
 - People think hard about *ridge operators* which let you jump around the space in better ways