

Artificial Intelligence

For HEDSPI Project

Lecturer 6 - Advanced search methods

Lecturers :

Dr. Le Thanh Huong

Dr. Tran Duc Khanh

Dr. Hai V. Pham

HUST

1


Outline

- Local beam search
- Game and search
- Alpha-beta pruning


2
2

Local beam search

- Like greedy search, but keep K states at all times:
 - Initially: k random states
 - Next: determine all successors of k states
 - If any of successors is goal \rightarrow finished
 - Else select k best from successors and repeat.



Greedy Search



Beam Search

3

Local beam search

- Major difference with random-restart search
 - Information is shared among k search threads: If one state generated good successor, but others did not \rightarrow “come here, the grass is greener!”
- Can suffer from lack of diversity.
 - Stochastic variant: choose k successors at proportionally to state success.
- The best choice in MANY practical settings

4

Games and search

- Why study games?
- Why is search a good idea?
- Major assumptions about games:
 - Only an agent's actions change the world
 - World is deterministic and accessible

5

Why study games?



May 1997
Deep Blue - Garry Kasparov
3.5 - 2.5

machines are better than humans in:
othello
humans are better than machines in:
go
here: perfect information zero-sum games

6

Why study games?

- Games are a form of *multi-agent environment*
 - What do other agents do and how do they affect our success?
 - Cooperative vs. competitive multi-agent environments.
 - Competitive multi-agent environments give rise to adversarial search a.k.a. *games*
- Why study games?
 - Fun; historically entertaining
 - Interesting subject of study because they are hard
 - Easy to represent and agents restricted to small number of actions

7

Relation of Games to Search

- Search – no adversary
 - Solution is (heuristic) method for finding goal
 - Heuristics and CSP techniques can find *optimal* solution
 - Evaluation function: estimate of cost from start to goal through given node
 - Examples: path planning, scheduling activities
- Games – adversary
 - Solution is strategy (strategy specifies move for every possible opponent reply).
 - Time limits force an *approximate* solution
 - Evaluation function: evaluate “goodness” of game position
 - Examples: chess, checkers, Othello, backgammon
- Ignoring computational complexity, games are a perfect application for a complete search.
- Of course, ignoring complexity is a bad idea, so games are a good place to study resource bounded searches.

8

Types of Games

	deterministic	chance
perfect information	chess, checkers, go, othello	backgammon monopoly
imperfect information	battleships, blind tictactoe	bridge, poker, scrabble nuclear war

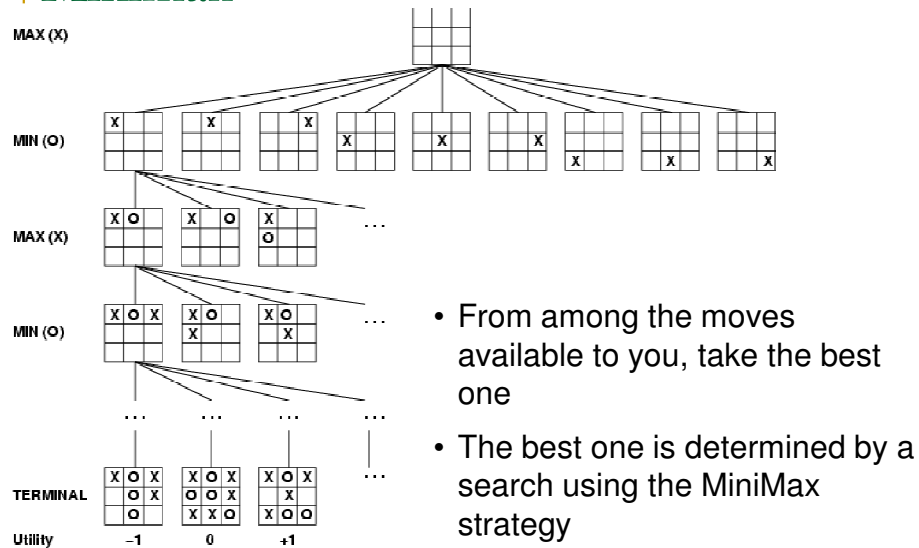
9

Minimax

- Two players: MAX and MIN
- MAX moves first and they take turns until the game is over. Winner gets award, loser gets penalty.
- Games as search:
 - **Initial state:** e.g. board configuration of chess
 - **Successor function:** list of (move,state) pairs specifying legal moves.
 - **Terminal test:** Is the game finished?
 - **Utility function:** Gives numerical value of terminal states.
 - E.g. win (+1), loose (-1) and draw (0) in tic-tac-toe
- MAX uses search tree to determine next move.
- Perfect play for deterministic games

10

Minimax



11

Optimal strategies

- MAX maximizes a function: find a move corresponding to max value
- MIN minimizes the same function: find a move corresponding to min value

At each step:

- If a state/node corresponds to a MAX move, the function value will be the maximum value of its childs
- If a state/node corresponds to a MIN move, the function value will be the minimum value of its childs

Given a game tree, the optimal strategy can be determined by using the minimax value of each node:

MINIMAX-VALUE(n) =

UTILITY(n)

If n is a terminal

$\max_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s)$


If n is a max node

$\min_{s \in \text{successors}(n)} \text{MINIMAX-VALUE}(s)$

If n is a min node

12

Minimax



13

Minimax algorithm

function MINIMAX-DECISION(*state*) *returns an action*

$v \leftarrow \text{MAX-VALUE}(\text{state})$

return the action in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return *v*

function MIN-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow \infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return *v*

14

Properties of minimax

- Complete? Yes (if tree is finite)
- Optimal? Yes (against an optimal opponent)
- Time complexity? $O(b^m)$
- Space complexity? $O(bm)$ (depth-first exploration)

- For chess, $b \approx 35$, $m \approx 100$ for "reasonable" games
→ exact solution completely infeasible

15

Problem of minimax search

- Number of games states is exponential to the number of moves.
 - Solution: Do not examine every node
- ⇒ Alpha-beta pruning:
- Remove branches that do not influence final decision
 - Revisit example ...


16

α - β pruning

- Alpha values: the best values achievable for MAX, hence the max value so far
- Beta values: the best values achievable for MIN, hence the min value so far
- At MIN level: compare result V of node to alpha value. If $V > \alpha$, pass value to parent node and BREAK
- At MAX level: compare result V of node to beta value. If $V < \beta$, pass value to parent node and BREAK

17


α - β pruning



18


α - β pruning example

Compare result V of node to β . If $V < \beta$, pass value to parent node and BREAK




19

α - β pruning example




20

α - β pruning example



21

α - β pruning example



22


Properties of α - β

- Pruning **does not** affect final result
- Entire sub-trees can be pruned.
- Good move ordering improves effectiveness of pruning. With "perfect ordering"
 - time complexity = $O(b^{m/2})$
 - **doubles** depth of search
 - Branching factor of \sqrt{b} !!
 - Alpha-beta pruning can look twice as far as minimax in the same amount of time
- Repeated states are again possible.
 - Store them in memory = transposition table
- A simple example of the value of reasoning about which computations are relevant (a form of **metareasoning**)

23

Why is it called α - β ?

- α is the value of the best (i.e., highest-value) choice found so far at any choice point along the path for *max*
- If v is worse than α , *max* will avoid it
 - prune that branch
- Define β similarly for *min*



24

The α - β algorithm

function ALPHA-BETA-SEARCH(*state*) *returns an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the action in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*, α , β) *returns a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** *v*

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return *v*

25

The α - β algorithm

function MIN-VALUE(*state*, α , β) *returns a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow +\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$

if $v \leq \alpha$ **then return** *v*

$\beta \leftarrow \text{MIN}(\beta, v)$

return *v*

26

Imperfect, real-time decisions

- Minimax and alpha-beta pruning require too much leafnode evaluations.
- May be impractical within a reasonable amount of time.
- Suppose we have 100 secs, explore 10^4 nodes/sec
→ 10^6 nodes per move
- Standard approach (SHANNON, 1950):
 - Cut off search earlier (replace TERMINAL-TEST by CUTOFF-TEST)
 - Apply heuristic evaluation function EVAL (replacing utility function of alpha-beta)

27

Cut-off search

- Change:
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
into:
 if CUTOFF-TEST(*state*,*depth*) **then return** EVAL(*state*)
- Introduces a fixed-depth limit *depth*
 - Is selected so that the amount of time will not exceed what the rules of the game allow.
- When cut-off occurs, the evaluation is performed.

28

Heuristic evaluation (EVAL)

- Idea: produce an estimate of the expected utility of the game from a given position.
- Requirements:
 - EVAL should order terminal-nodes in the same way as UTILITY.
 - Computation may not take too long.
 - For non-terminal states the EVAL should be strongly correlated with the actual chance of winning.
- Example:

Expected value $e(p)$ for each state p :


$$E(p) = (\text{\# open rows, columns, diagonals for MAX}) - (\text{\# open rows, columns, diagonals for MIN})$$
 - MAX moves all lines that don't have o; MIN moves all lines that don't have x

29

Expected value $e(p)$ for each state p :

$$E(p) = (\text{\# open rows, columns, diagonals for MAX}) - (\text{\# open rows, columns, diagonals for MIN})$$


MAX moves all lines that don't have o; MIN moves all lines that don't have x



→ A kind of depth-first search


30 30

Evaluation function example



Black to move

White slightly better



White to move

Black winning

- For chess, typically **linear** weighted sum of **features**
$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$
- e.g., $w_1 = 9$ with
 $f_1(s) = (\text{number of white queens}) - (\text{number of black queens}), \text{ etc.}$

31

Chess complexity

- PC can search 200 millions nodes/3min.
- Branching factor: ~35
 - $35^5 \sim 50$ millions
 - if use minimax, could look ahead **5 plies**, defeated by average player, planning 6-8 plies.
- Does it work in practice?
 - 4-ply \approx human novice \rightarrow hopeless chess player
 - 8-ply \approx typical PC, human master
 - 12-ply \approx Deep Blue, Kasparov
- To reach grandmaster level, needs a better *extensively tuned evaluation* and a *large database of optimal opening and ending* of the game

32


Deterministic games in practice

- Checkers: Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a precomputed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 444 billion positions.
- Chess: Deep Blue defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.
- Othello: human champions refuse to compete against computers, who are too good.
- Go: human champions refuse to compete against computers, who are too bad. In go, $b > 300$, so most programs use pattern knowledge bases to suggest plausible moves.

33

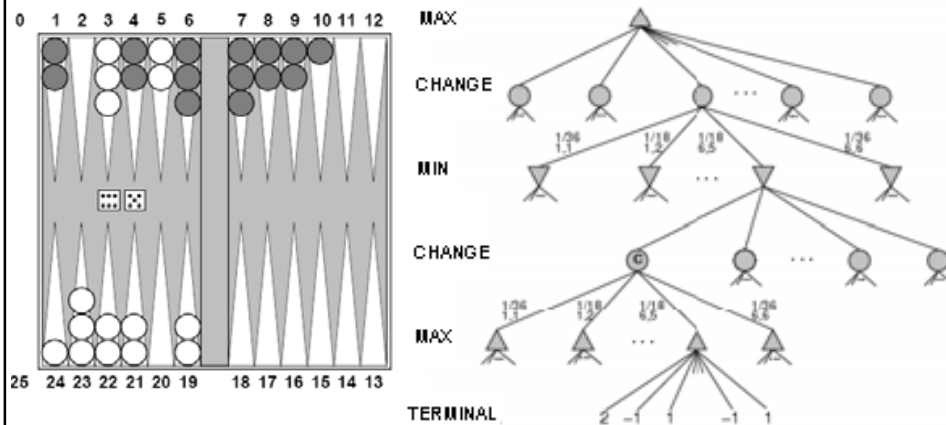
Nondeterministic games

- Chance introduces by dice, card-shuffling, coin-flipping...
- Example with coin-flipping:



34

Backgammon



Possible moves: (5-10,5-11), (5-11,19-24), (5-10,10-16) and (5-11,11-16)

35

Expected minimax value

```

...
if state is a MAX node then
    return the highest EXPECTIMINIMAX-VALUE of SUCCESSORS(state)
if state is a MIN node then
    return the lowest EXPECTIMINIMAX-VALUE of SUCCESSORS(state)
if state is a chance node then
    return average of EXPECTIMINIMAX-VALUE of SUCCESSORS(state)
    
```

EXPECTED-MINIMAX-VALUE(n)=

$UTILITY(n)$ If n is a terminal
 $\max_{s \in \text{successors}(n)} \text{EXPECTEDMINIMAX}(s)$ If n is a max node
 $\min_{s \in \text{successors}(n)} \text{EXPECTEDMINIMAX}(s)$ If n is a min node
 $\sum_{s \in \text{successors}(n)} P(s) \cdot \text{EXPECTEDMINIMAX}(s)$ If n is a chance node

$P(s)$ is probability of s occurrence

36

Games of imperfect information

- E.g., card games, where opponent's initial cards are unknown
- Typically we can calculate a probability for each possible deal
- Seems just like having one big dice roll at the beginning of the game
- Idea: compute the minimax value of each action in each deal, then choose the action with highest expected value over all deals
- Special case: if an action is optimal for all deals, it's optimal.
- GIB, current best bridge program, approximates this idea by
 - generating 100 deals consistent with bidding information
 - picking the action that wins most tricks on average